

▼ CROP RECOMMENDATION SYSTEM

```
# Importing libraries

from __future__ import print_function
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn import tree
import warnings
warnings.filterwarnings('ignore')

from google.colab import files
uploaded = files.upload()



Choose Files



No file chosen



Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.



Saving Crop_recommendation.csv to Crop_recommendation.csv



df = pd.read_csv("Crop_recommendation.csv")

df.head()



|   | N  | P  | K  | temperature | humidity  | ph       | rainfall   | label |
|---|----|----|----|-------------|-----------|----------|------------|-------|
| 0 | 90 | 42 | 43 | 20.879744   | 82.002744 | 6.502985 | 202.935536 | rice  |
| 1 | 85 | 58 | 41 | 21.770462   | 80.319644 | 7.038096 | 226.655537 | rice  |
| 2 | 60 | 55 | 44 | 23.004459   | 82.320763 | 7.840207 | 263.964248 | rice  |
| 3 | 74 | 35 | 40 | 26.491096   | 80.158363 | 6.980401 | 242.864034 | rice  |
| 4 | 78 | 42 | 42 | 20.130175   | 81.604873 | 7.628473 | 262.717340 | rice  |



df.tail()



|      | N   | P  | K  | temperature | humidity  | ph       | rainfall   | label  |
|------|-----|----|----|-------------|-----------|----------|------------|--------|
| 2195 | 107 | 34 | 32 | 26.774637   | 66.413269 | 6.780064 | 177.774507 | coffee |
| 2196 | 99  | 15 | 27 | 27.417112   | 56.636362 | 6.086922 | 127.924610 | coffee |
| 2197 | 118 | 33 | 30 | 24.131797   | 67.225123 | 6.362608 | 173.322839 | coffee |
| 2198 | 117 | 32 | 34 | 26.272418   | 52.127394 | 6.758793 | 127.175293 | coffee |
| 2199 | 104 | 18 | 30 | 23.603016   | 60.396475 | 6.779833 | 140.937041 | coffee |



df.size

17600

df.shape

(2200, 8)

df.columns

Index(['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall', 'label'], dtype='object')

df['label'].unique()

array(['rice', 'maize', 'chickpea', 'kidneybeans', 'pigeonpeas', 'mothbeans', 'mungbean', 'blackgram', 'lentil', 'pomegranate', 'banana', 'mango', 'grapes', 'watermelon', 'muskmelon', 'apple', 'orange', 'papaya', 'coconut', 'cotton', 'jute', 'coffee'], dtype=object)

df.dtypes
```

```

N          int64
P          int64
K          int64
temperature float64
humidity    float64
ph          float64
rainfall    float64
label      object
dtype: object

```

```
df['label'].value_counts()
```


```

rice      100
maize     100
jute      100
cotton    100
coconut   100
papaya    100
orange    100
apple     100
muskmelon 100
watermelon 100
grapes    100
mango     100
banana    100
pomegranate 100
lentil    100
blackgram 100
mungbean  100
mothbeans 100
pigeonpeas 100
kidneybeans 100
chickpea  100
coffee    100
Name: label, dtype: int64

```

```
sns.heatmap(df.corr(),annot=True)
```

```
<Axes: >
```



▼ Separating features and target label

```

features = df[['N', 'P','K','temperature', 'humidity', 'ph', 'rainfall']]
target = df['label']
labels = df['label']

# Initializing empty lists to append all model's name and corresponding name
acc = []
model = []

```

```
# Splitting into train and test data

from sklearn.model_selection import train_test_split
Xtrain, Xtest, Ytrain, Ytest = train_test_split(features,target,test_size = 0.2,random_state =2)
```

▼ Decision Tree

```
from sklearn.tree import DecisionTreeClassifier

DecisionTree = DecisionTreeClassifier(criterion="entropy",random_state=2,max_depth=5)

DecisionTree.fit(Xtrain,Ytrain)

predicted_values = DecisionTree.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Decision Tree')
print("DecisionTrees's Accuracy is: ", x*100)

print(classification_report(Ytest,predicted_values))
```

```
DecisionTrees's Accuracy is: 90.0
              precision    recall  f1-score   support

   apple         1.00        1.00        1.00         13
  banana         1.00        1.00        1.00         17
blackgram         0.59        1.00        0.74         16
  chickpea         1.00        1.00        1.00         21
   coconut         0.91        1.00        0.95         21
   coffee         1.00        1.00        1.00         22
   cotton         1.00        1.00        1.00         20
   grapes         1.00        1.00        1.00         18
     jute         0.74        0.93        0.83         28
kidneybeans         0.00        0.00        0.00         14
   lentil         0.68        1.00        0.81         23
     maize         1.00        1.00        1.00         21
     mango         1.00        1.00        1.00         26
  mothbeans         0.00        0.00        0.00         19
   mungbean         1.00        1.00        1.00         24
 muskmelon         1.00        1.00        1.00         23
   orange         1.00        1.00        1.00         29
   papaya         1.00        0.84        0.91         19
pigeonpeas         0.62        1.00        0.77         18
pomegranate         1.00        1.00        1.00         17
     rice         1.00        0.62        0.77         16
watermelon         1.00        1.00        1.00         15

 accuracy                   0.90         440
  macro avg         0.84        0.88        0.85         440
 weighted avg         0.86        0.90        0.87         440
```

```
from sklearn.model_selection import cross_val_score

# Cross validation score (Decision Tree)
score = cross_val_score(DecisionTree, features, target,cv=5)

score

array([0.93636364, 0.90909091, 0.91818182, 0.87045455, 0.93636364])
```

▼ Saving trained Decision Tree model

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
DT_pkl_filename = 'DecisionTree.pkl'
# Open the file to save as pkl file
DT_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(DecisionTree, DT_Model_pkl)
# Close the pickle instances
DT_Model_pkl.close()
```

▼ Guassian Naive Bayes

```

from sklearn.naive_bayes import GaussianNB

NaiveBayes = GaussianNB()

NaiveBayes.fit(Xtrain,Ytrain)

predicted_values = NaiveBayes.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Naive Bayes')
print("Naive Bayes's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))


```

	Naive Bayes's Accuracy is:	0.990909090909091		
	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	1.00	1.00	1.00	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	18
jute	0.88	1.00	0.93	28
kidneybeans	1.00	1.00	1.00	14
lentil	1.00	1.00	1.00	23
maize	1.00	1.00	1.00	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	1.00	1.00	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	1.00	0.75	0.86	16
watermelon	1.00	1.00	1.00	15
accuracy			0.99	440
macro avg	0.99	0.99	0.99	440
weighted avg	0.99	0.99	0.99	440

```

# Cross validation score (NaiveBayes)
score = cross_val_score(NaiveBayes,features,target,cv=5)
score

array([0.99772727, 0.99545455, 0.99545455, 0.99545455, 0.99090909])

```

▼ Saving trained Guassian Naive Bayes model

```

import pickle
# Dump the trained Naive Bayes classifier with Pickle
NB_pkl_filename = 'NBClassifier.pkl'
# Open the file to save as pkl file
NB_Model_pkl = open(NB_pkl_filename, 'wb')
pickle.dump(NaiveBayes, NB_Model_pkl)
# Close the pickle instances
NB_Model_pkl.close()

```

▼ Support Vector Machine (SVM)

```

from sklearn.svm import SVC

SVM = SVC(gamma='auto')

SVM.fit(Xtrain,Ytrain)

predicted_values = SVM.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('SVM')
print("SVM's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))

```

```
SVM's Accuracy is: 0.10681818181818181
precision    recall  f1-score   support

   apple      1.00      0.23      0.38        13
  banana      1.00      0.24      0.38        17
blackgram      1.00      0.19      0.32        16
  chickpea      1.00      0.05      0.09        21
   coconut      1.00      0.05      0.09        21
    coffee      0.00      0.00      0.00        22
   cotton      1.00      0.05      0.10        20
   grapes      1.00      0.06      0.11        18
     jute      1.00      0.07      0.13        28
kidneybeans    0.03      1.00      0.07        14
   lentil      0.00      0.00      0.00        23
    maize      0.00      0.00      0.00        21
    mango      0.00      0.00      0.00        26
  mothbeans    0.00      0.00      0.00        19
  mungbean      1.00      0.12      0.22        24
 muskmelon      1.00      0.30      0.47        23
   orange      1.00      0.03      0.07        29
   papaya      1.00      0.05      0.10        19
pigeonpeas     0.00      0.00      0.00        18
pomegranate     1.00      0.12      0.21        17
     rice      0.50      0.06      0.11        16
 watermelon     1.00      0.13      0.24        15

 accuracy                    0.11       440
  macro avg      0.66      0.13      0.14       440
 weighted avg     0.66      0.11      0.13       440
```

```
# Cross validation score (SVM)
score = cross_val_score(SVM, features, target, cv=5)
score

array([0.27727273, 0.28863636, 0.29090909, 0.275      , 0.26818182])
```

▼ Logistic Regression

```
from sklearn.linear_model import LogisticRegression

LogReg = LogisticRegression(random_state=2)

LogReg.fit(Xtrain, Ytrain)

predicted_values = LogReg.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Logistic Regression')
print("Logistic Regression's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))

Logistic Regression's Accuracy is: 0.9522727272727273
precision    recall  f1-score   support

   apple      1.00      1.00      1.00        13
  banana      1.00      1.00      1.00        17
blackgram     0.86      0.75      0.80        16
  chickpea      1.00      1.00      1.00        21
   coconut      1.00      1.00      1.00        21
    coffee      1.00      1.00      1.00        22
   cotton     0.86      0.90      0.88        20
   grapes      1.00      1.00      1.00        18
     jute     0.84      0.93      0.88        28
kidneybeans    1.00      1.00      1.00        14
   lentil     0.88      1.00      0.94        23
    maize     0.90      0.86      0.88        21
    mango     0.96      1.00      0.98        26
  mothbeans     0.84      0.84      0.84        19
  mungbean      1.00      0.96      0.98        24
 muskmelon      1.00      1.00      1.00        23
   orange      1.00      1.00      1.00        29
   papaya      1.00      0.95      0.97        19
pigeonpeas     1.00      1.00      1.00        18
pomegranate     1.00      1.00      1.00        17
     rice     0.85      0.69      0.76        16
 watermelon     1.00      1.00      1.00        15

 accuracy                    0.95       440
  macro avg      0.95      0.95      0.95       440
```

```
weighted avg      0.95      0.95      0.95      440
```

```
# Cross validation score (Logistic Regression)
score = cross_val_score(LogReg, features, target, cv=5)
score

array([0.95      , 0.96590909, 0.94772727, 0.96590909, 0.94318182])
```

▼ Saving trained Logistic Regression model

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
LR_pkl_filename = 'LogisticRegression.pkl'
# Open the file to save as pkl file
LR_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(LogReg, LR_Model_pkl)
# Close the pickle instances
LR_Model_pkl.close()
```

▼ Random Forest

```
from sklearn.ensemble import RandomForestClassifier

RF = RandomForestClassifier(n_estimators=20, random_state=0)
RF.fit(Xtrain, Ytrain)

predicted_values = RF.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('RF')
print("RF's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))
```

```
RF's Accuracy is: 0.990909090909091
      precision    recall  f1-score   support

   apple         1.00      1.00      1.00        13
  banana         1.00      1.00      1.00        17
blackgram         0.94      1.00      0.97        16
  chickpea         1.00      1.00      1.00        21
   coconut         1.00      1.00      1.00        21
   coffee         1.00      1.00      1.00        22
   cotton         1.00      1.00      1.00        20
   grapes         1.00      1.00      1.00        18
     jute         0.90      1.00      0.95        28
kidneybeans         1.00      1.00      1.00        14
   lentil         1.00      1.00      1.00        23
   maize         1.00      1.00      1.00        21
   mango         1.00      1.00      1.00        26
  mothbeans         1.00      0.95      0.97        19
   mungbean         1.00      1.00      1.00        24
 muskmelon         1.00      1.00      1.00        23
   orange         1.00      1.00      1.00        29
   papaya         1.00      1.00      1.00        19
pigeonpeas         1.00      1.00      1.00        18
pomegranate         1.00      1.00      1.00        17
     rice         1.00      0.81      0.90        16
watermelon         1.00      1.00      1.00        15

 accuracy                   0.99        440
  macro avg              0.99      0.99      0.99        440
 weighted avg              0.99      0.99      0.99        440
```

```
# Cross validation score (Random Forest)
score = cross_val_score(RF, features, target, cv=5)
score

array([0.99772727, 0.99545455, 0.99772727, 0.99318182, 0.98863636])
```


▼ Saving trained Random Forest model

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
RF_pkl_filename = 'RandomForest.pkl'
# Open the file to save as pkl file
RF_Model_pkl = open(RF_pkl_filename, 'wb')
pickle.dump(RF, RF_Model_pkl)
# Close the pickle instances
RF_Model_pkl.close()
```

▼ Accuracy Comparison

```
plt.figure(figsize=[10,5],dpi = 100)
plt.title('Accuracy Comparison')
plt.xlabel('Accuracy')
plt.ylabel('Algorithm')
sns.barplot(x = acc,y = model,palette='dark')
```

```
<Axes: title={'center': 'Accuracy Comparison'}, xlabel='Accuracy', ylabel='Algorithm'>
```



```
accuracy_models = dict(zip(model, acc))
for k, v in accuracy_models.items():
    print (k, '-->', v)

Decision Tree --> 0.9
Naive Bayes --> 0.990909090909091
SVM --> 0.106818181818181
Logistic Regression --> 0.9522727272727273
RF --> 0.990909090909091
```

▼ Making a prediction

```
data = np.array([[104,18, 30, 23.603016, 60.3, 6.7, 140.91]])
prediction = RF.predict(data)
print(prediction)
```

```
['coffee']
```

```
data = np.array([[83, 45, 60, 28, 70.3, 7.0, 150.9]])
prediction = RF.predict(data)
print(prediction)
```

```
['jute']
```

